Grimjack at Touché 2022

Advanced IR, Winter Semester 2021/22

Johannes Huck Jan Heinrich Reimer

Martin Luther University Halle-Wittenberg

February 7, 2022

Task at hand

- Task 2 of Touché: Argument Retrieval
- Argument Retrieval for Comparative Questions
- Task: Retrieve relevant passages to answer comparative questions and detect their stance w.r.t the objects
- Data: > 1 million text passages

https://mobile.twitter.com/webis_de/status/1468529926026534913?cxt=HHwWgoC97fyLouEoAAAA

General approach

- Programmed in Python
  - Easy to use
  - High readability
  - Many IR libabries available
- Three modules: Search, Run file and Evaluate
- Pipeline consists of
  - Query-Expander and Query-Combiner
  - Initial Retrieval
  - Argument quality and stance tagging
  - Reranking
- Indexing and initial retrieval via pyserini [Lin+21]

Pipeline

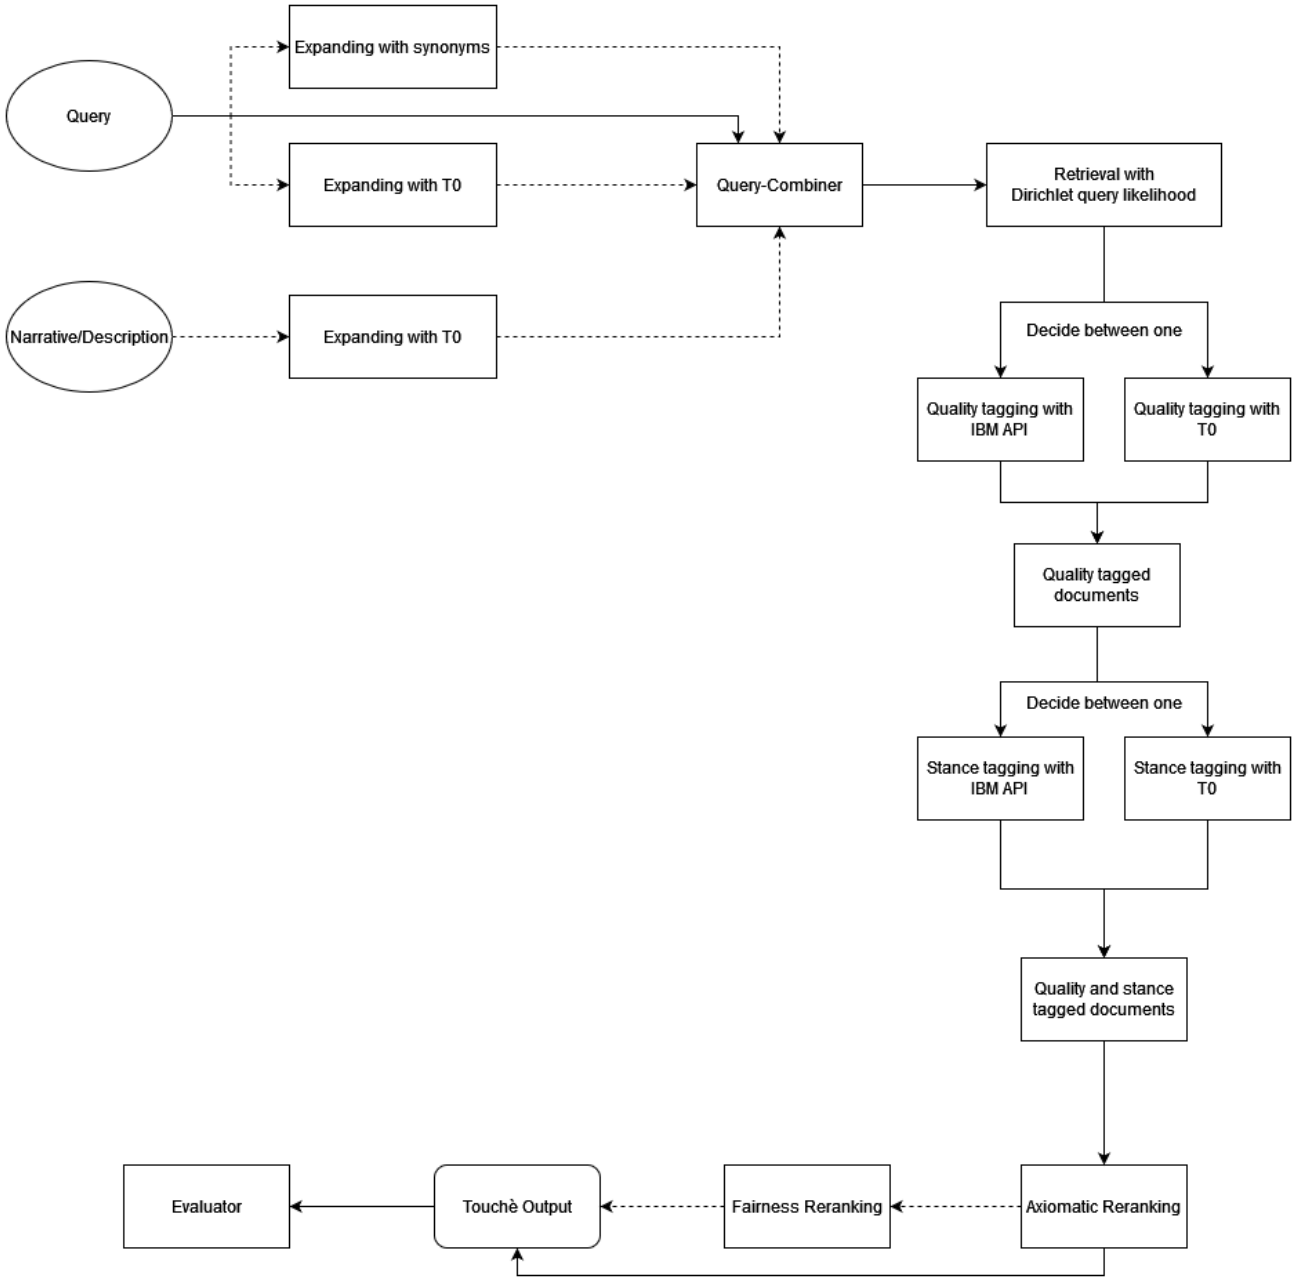
Query
Expanding with synonyms
Expanding with T0
Narrative/Description
Expanding with T0
Query-Combiner
Retrieval with Dirichlet query likelihood
Decide between one
Quality tagging with IBM API
Quality tagging with T0
Quality tagged documents
Decide between one
Stance tagging with IBM API
Stance tagging with T0
Quality and stance tagged documents
Axiomatic Reranking
Fairness Reranking
Touchè Output
Evaluator

Query-Expander and Query-Combiner

- Expanding queries with synonyms of comparative objects
- Two Different approaches
  - Based on embeddings with glove
  - Based on language model T0 [San+21]
  - We ask "What are synonyms of the word <token> ?"
- With T0 also new queries from narrative and description
- We ask "<text> Extract a natural search query from this description."
- Combining all new queries with OR
- Retrieving ranked list of passages with this new query

Argument quality tagging

- Extracting arguments with TARGER [Che+19]
- For each argument we want to know the quality w.r.t. the topic
- Two different approaches
  - Based IBM Debater API [Tol+19]
  - Based on T0
  - We ask "<sentence> How would you rate the readability and consistency in this sentence? very good, good, bad, very bad"
- IBM Debater API returns a score between 0 and 1
- 0 means lowest quality and 1 highest quality

Example

Arg: Cars should only provide assisted driving, not complete autonomy
Topic: We should further explore the development of autonomous vehicles
Score: 0.7256

Argument stance tagging

- Next we want to know the stance w.r.t. the topic
- Two different approaches
  - Based on IBM Debater API [Bar+17]
  - Based on T0
  - We ask "<sentence> Is this sentence pro/against <comparative_object>? yes or no"
- It is also possible to expand with sentiments
- Both approaches only work for single target stance
- Calculating the multi target stance
  - Calculate the difference between objects
  - Use a threshold
  - Convert T0s output into a numerical representation

Axiomatic Reranking

- Compute preferences between documents ($\triangleq$ axioms)
- Multiple axioms vote against the original ranking
- Rerank with KwikSort [Hag+16]

Argumentative Axioms

ArgUC Prefer more argumentative units [Bon+18]

QTArg Prefer more query terms in argumentative units [Bon+18]

QTPArg Prefer earlier query terms in argumentative units [Bon+18]

aSL Prefer sentences with 12–20 words [Bon+21]

CompArg Prefer more comparative objects in argumentative units

CompPArg Prefer earlier comparative objects in argumentative units

ArgQ Prefer higher argument quality

Fairness Reranking

- Idea: prefer subjective arguments over neutral arguments
  but guarantee fair exposure for each stance (pro/con)
- Alternating stance
  - Three filtered lists by stance: first, second, neutral/other
  - Alternately select from first/second list
  - Fallback to neutral list if first/second list is empty
- Balanced top-k stance
  - Count number of documents pro first or pro second in top-k ranking
  - If difference > 1:
    Move last pro first document from top-k ranking
    after the first pro second document after top-k ranking

Final Remarks

- Approach is very flexible
- We investigate influence of components w.r.t the retrieval score
- Stance classification may be better with Roberta approach
- We cannot distinguish between neutral and no stance
- We investigate how reranking influences the retrieval score
- T0 solves a lot of IR tasks
- Is it possible to only use T0 for retrieval?

Thank you!

References

- Bar-Haim, Roy et al. (2017). "Stance Classification of Context-Dependent Claims". In: *EACL*.
- Bondarenko, Alexander et al. (Nov. 2018). "Webis at TREC 2018: Common Core Track". In: *27th International Text Retrieval Conference (TREC 2018)*. Ed. by Ellen M. Voorhees et al. NIST Special Publication. National Institute of Standards and Technology (NIST).
- Bondarenko, Alexander et al. (2021). "Axiomatic Re-Ranking for Argument Retrieval". In:
- Chernodub, Artem N. et al. (2019). "TARGER: Neural Argument Mining at Your Fingertips". In: *Proceedings of the 57th Conference of the Association for Computational Linguistics, ACL 2019, Florence, Italy, July 28 - August 2, 2019, Volume 3: System Demonstrations*. Ed. by Marta R. Costa-jussà et al. Association for Computational Linguistics, pp. 195–200.
- Hagen, Matthias et al. (Oct. 2016). "Axiomatic Result Re-Ranking". In: *25th ACM International Conference on Information and Knowledge Management (CIKM 2016)*. ACM, pp. 721–730.
- Lin, Jimmy J. et al. (2021). "Pyserini: An Easy-to-Use Python Toolkit to Support Replicable IR Research with Sparse and Dense Representations". In: *ArXiv* abs/2102.10073.
- Sanh, Victor et al. (2021). "Multitask Prompted Training Enables Zero-Shot Task Generalization". In: *CoRR* abs/2110.08207. arXiv: `2110.08207`.

References (cont.)

- Toledo, Assaf et al. (2019). "Automatic Argument Quality Assessment - New Datasets and Methods". In: *EMNLP*.